

In [1]:

```
#Assignment 2: Implementing Feedforward neural networks with Keras and TensorFlow
from sklearn.preprocessing import LabelBinarize
from sklearn.metrics import classification_report
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.datasets import mnist
from tensorflow.keras import backend as K
import matplotlib.pyplot as plt
import numpy as np
```

In [2]:

```
((X_train, Y_train), (X_test, Y_test)) = mnist.load_data()
X_train = X_train.reshape((X_train.shape[0], 28 * 28 * 1))
X_test = X_test.reshape((X_test.shape[0], 28 * 28 * 1))
X_train = X_train.astype("float32") / 255.0
X_test = X_test.astype("float32") / 255.0
```

In [3]:

```
lb = LabelBinarizer()
Y_train = lb.fit_transform(Y_train)
Y_test = lb.transform(Y_test)
```

In [4]:

```
model = Sequential()
model.add(Dense(128, input_shape=(784,), activation="sigmoid"))
model.add(Dense(64, activation="sigmoid"))
model.add(Dense(10, activation="softmax"))
```

```
In [8]: sgd = SGD(0.01)
epochs=10
model.compile(loss="categorical_crossentropy", optimizer=sgd,metrics=["accuracy"]
H = model.fit(X_train, Y_train, validation_data=(X_test, Y_test),epochs=epochs,b
```

```
Epoch 1/10
469/469 [=====] - 4s 7ms/step - loss: 2.2846 - accurac
y: 0.1874 - val_loss: 2.2474 - val_accuracy: 0.2559
Epoch 2/10
469/469 [=====] - 3s 6ms/step - loss: 2.2204 - accurac
y: 0.3637 - val_loss: 2.1847 - val_accuracy: 0.4210
Epoch 3/10
469/469 [=====] - 3s 6ms/step - loss: 2.1479 - accurac
y: 0.4922 - val_loss: 2.0960 - val_accuracy: 0.5758
Epoch 4/10
469/469 [=====] - 3s 6ms/step - loss: 2.0428 - accurac
y: 0.5624 - val_loss: 1.9677 - val_accuracy: 0.6247
Epoch 5/10
469/469 [=====] - 3s 6ms/step - loss: 1.8940 - accurac
y: 0.6043 - val_loss: 1.7939 - val_accuracy: 0.6364
Epoch 6/10
469/469 [=====] - 3s 6ms/step - loss: 1.7091 - accurac
y: 0.6344 - val_loss: 1.5973 - val_accuracy: 0.6702
Epoch 7/10
469/469 [=====] - 3s 6ms/step - loss: 1.5184 - accurac
y: 0.6687 - val_loss: 1.4123 - val_accuracy: 0.6941
Epoch 8/10
469/469 [=====] - 3s 6ms/step - loss: 1.3480 - accurac
y: 0.6986 - val_loss: 1.2550 - val_accuracy: 0.7204
Epoch 9/10
469/469 [=====] - 3s 6ms/step - loss: 1.2052 - accurac
y: 0.7268 - val_loss: 1.1259 - val_accuracy: 0.7414
Epoch 10/10
469/469 [=====] - 3s 6ms/step - loss: 1.0875 - accurac
y: 0.7487 - val_loss: 1.0191 - val_accuracy: 0.7692
```

```
In [10]: predictions = model.predict(X_test, batch_size=128)
print(classification_report(Y_test.argmax(axis=1),predictions.argmax(axis=1),tar
```


```
79/79 [=====] - 0s 4ms/step
              precision    recall  f1-score   support

     0       0.82         0.97         0.89         980
     1       0.77         0.99         0.87        1135
     2       0.84         0.70         0.76        1032
     3       0.67         0.83         0.74        1010
     4       0.71         0.88         0.78         982
     5       0.78         0.41         0.54         892
     6       0.81         0.87         0.84         958
     7       0.76         0.88         0.81        1028
     8       0.79         0.54         0.64         974
     9       0.79         0.56         0.65        1009

 accuracy         0.77         0.77         0.75        10000
 macro avg       0.77         0.76         0.75        10000
 weighted avg    0.77         0.77         0.76        10000
```

```
In [11]: plt.style.use("ggplot")
plt.figure()
plt.plot(np.arange(0, epochs), H.history["loss"], label="train_loss")
plt.plot(np.arange(0, epochs), H.history["val_loss"], label="val_loss")
plt.plot(np.arange(0, epochs), H.history["accuracy"], label="train_acc")
plt.plot(np.arange(0, epochs), H.history["val_accuracy"], label="val_acc")
plt.title("Training Loss and Accuracy")
plt.xlabel("Epoch #")
plt.ylabel("Loss/Accuracy")
plt.legend()
```

Out[11]: <matplotlib.legend.Legend at 0x1ee51a68d90>



In []: